FlashPASTA: taking Fourier augmentation off the CPU critical path

Maryam Saif[1], **Esraa Saif**[2]
[1]SDAIA · [2]IBM
maryam.h.saif@gmail.com · esraa.saif@gmail.com

Saudi HPC/AI 2026
KAUST · 26–28 October
Methods that improve computational efficiency

Code & data

42×
FlashPASTA on 1 GPU vs official code on 32 CPU cores

77% → 0.2%
of training time the GPU sits waiting for data (8 CPU workers)

4.2×
faster YOLO training on Sim10k with the same 8 workers

3×10^{-6}
largest pixel difference from the official output, same noise

The problem

Synthetic frames come with free labels, but detectors trained on them lose accuracy on real photos. **PASTA** [1] closes part of the gap by randomly perturbing each image's Fourier amplitude spectrum, more strongly at high frequencies.



Sim10k frame [2]



After PASTA (α=3, β=0.25, k=2)

The official code runs per image inside CPU data-loader workers, with float64 FFTs and the noise mask rebuilt on every call: **0.7 images/s per core** at 1914×1052. At 640 px it slows a worker from **77 to 5.9 images/s**.

FlashPASTA

Official



FlashPASTA



- **Batched cuFFT on the GPU** after the batch is transferred
- **Noise mask cached** per image size
- **F·ε replaces $|F|\cdot\varepsilon\cdot e^{i\varphi}$**: no abs, angle or exp
- **Half-size real FFT**, exact because only the real part is kept:

```
Re IFFT(F·ε) = IRFFT(F·(ε(f)+ε(−f))/2)
```

Same output as the official code

Same noise: largest pixel difference (float32)	**3.3×10^{-6}**
8-bit pixels that change (by 1 level)	**0.010%**
Sampled noise spread vs official	**0.98–1.02×**
Pixel means that differ (|z| > 3)	**0 / 9,216**

1 · PASTA alone: 42× faster on one GPU

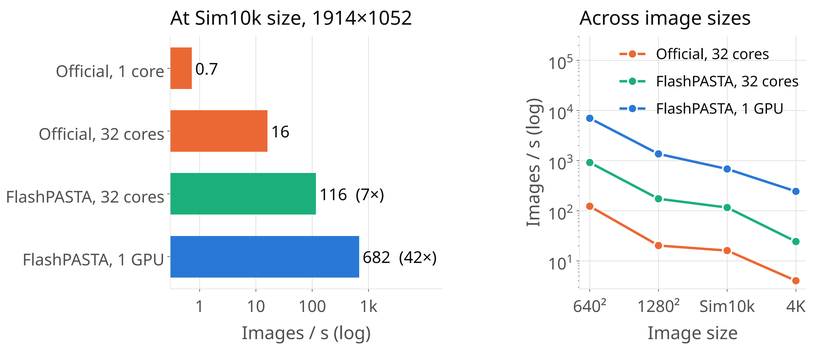


Figure 1. Augmentation only, no data loading. RTX A5000 vs Threadripper PRO 3975WX (32 cores); CPU runs use single-threaded processes, like data-loader workers.

2 · Data loading: GPU PASTA is nearly free

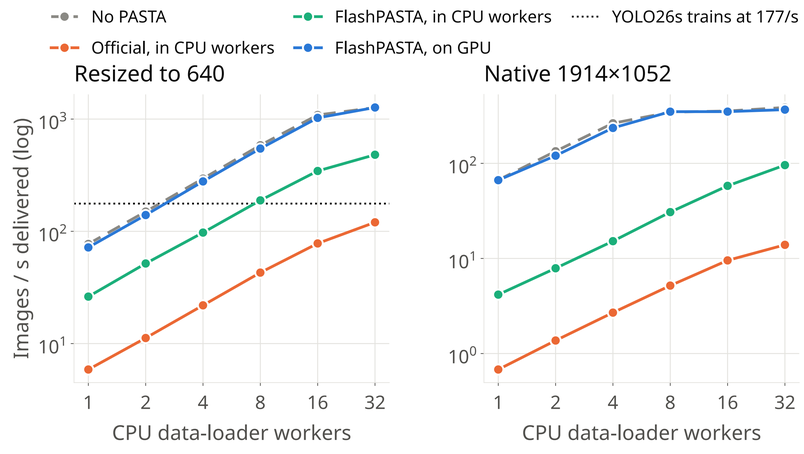


Figure 2. Images/s a PyTorch DataLoader delivers on Sim10k (JPEG decode → letterbox → PASTA) vs CPU workers.

Takeaways

- FlashPASTA on one GPU is 42× faster than the official code on 32 CPU cores, with the same output up to floating-point rounding.
- It no longer eats the CPU-core budget, so a node feeds its GPU with a few workers.
- Restructuring (6.4× on one core) and the real-FFT identity (1.4× on GPU) apply to any Fourier-domain augmentation.

3 · YOLO training: 4.2× faster epochs

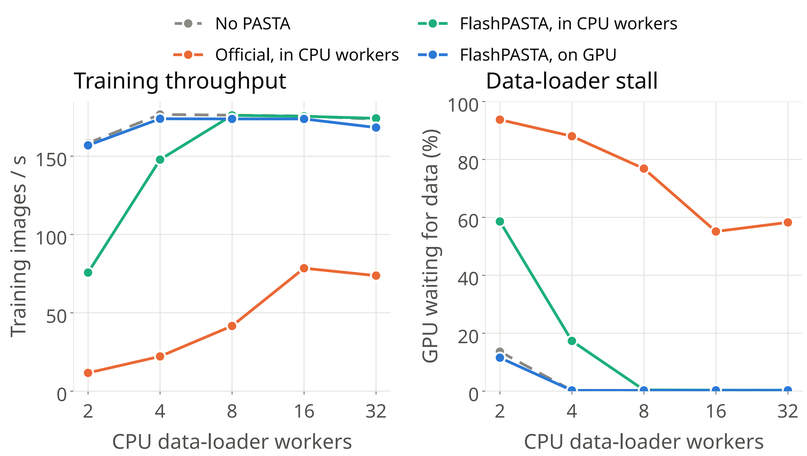


Figure 3. YOLO26s, 640 px, batch 16, one RTX A5000, mixed precision; mosaic off so every mode sees identical inputs.

8 CPU workers, 640 px	Official, CPU	FlashPASTA, GPU
Training throughput	42 img/s	**174 img/s**
GPU waiting for data	77%	**0.2%**
One Sim10k epoch (10k images)	4.0 min	**1.0 min**
100 epochs (projected)	6.7 h	**1.6 h**

The official version never feeds the GPU's 177 img/s, not even with 16 workers (78 img/s). At 1280 px the gap widens to **6.4×**.

4 · Red Sea corals: 3.6× faster training

Real underwater images, Marjan dataset [4] (9,498 train), same setup.



Healthy



Bleached



Dead

8 CPU workers, 640 px	Official, CPU	FlashPASTA, GPU
GPU waiting for data	74%	**0.2%**
One epoch (9,498 images)	3.4 min	**0.9 min**

[1] P. Chattopadhyay, K. Sarangmath, V. Vijaykumar, J. Hoffman. PASTA: Proportional Amplitude Spectrum Training Augmentation for Syn-to-Real Domain Generalization. ICCV 2023. [2] M. Johnson-Roberson et al. Driving in the Matrix. ICRA 2017. Sim10k, doi:10.7302/e1f1-3d97, CC BY-NC 4.0. [3] Ultralytics YOLO26. [4] A. Aldhahri, E. Saif, H. Ali, M. Alsayed, F. Alshareef. Harnessing Computer Vision and Deep Learning to Monitor Coral Reef Health. ETASR 15(4), 2025, doi:10.48084/etasr.10324; Marjan dataset, Roboflow Universe, CC BY 4.0.

github.com/Maryyyyyyyam142/flash-pasta